



Where we are • what's new • today's plan

~19 min • recap • AI in the news • opening quiz • the guest • A1 & A2 • today's plan

BY THE END OF THIS CAPSULE YOU WILL BE ABLE TO

Know where today sits in the course, what changed in the AI market this month, and exactly what we will decide — and price — over the next three hours.

Where we were — and what today does with it

Four classes in one line: the map, the stack around the model, inside the model, then the loop. Today: which of it is worth doing, and what it costs.

#	What Class 4 left you with	What Class 5 does with it
1	thought → action → observation	C2 prices it — the observation is re-sent every turn
2	errors compound across steps	C2's failure term — cost per SUCCESSFUL task
3	cost grows faster than the turns	C2 puts the formula up: $B \times T + \frac{1}{2}DT^2$
4	cost per task, not price per token	the spine of C2 — and why the ranking inverts
5	the first write is the governance cliff	C3 — Ant's internal fixes made an external one
6	readiness is set by the weakest link	C1's five data-readiness tests find that link

Takeaway Class 4 priced one agent. Class 5 decides which use cases deserve one at all.

AI in the news — the price fell again

THIS MONTH

Two items from this month, and together they are the argument of Capsule 2.

The frontier got cheaper — again, on 12 August. Grok 4.6 shipped at **\$2 in / \$6 out per million tokens**. The frontier tier you price in C2 is \$5 and \$25.

Why today: Force 1 — a fixed capability getting about 10× cheaper a year — is not history in a reading. Our own worked example is three weeks old and already the expensive option.

And the best model on the market could not sell. The Ramp AI Index (12 Aug) reports the top-performing model, at roughly \$10 per million tokens, took **6% of tokens and 11.4% of spend** on its platform in the month after launch. Buyers found a ceiling.


Why today: C2 spends a slide on this: “we used the best model” is not an argument, it is a cost. The market got there before we taught it.

Sources: xAI pricing, effective 12 Aug 2026 • Ramp AI Index, 12 Aug 2026 (US card and bill-pay data). Both checked 30 Aug 2026.

Takeaway Price per token keeps falling — and buyers still refuse to pay for the top of the range.

AI in the news — the spend has split in two

Median monthly AI spend per employee, US businesses, July 2026 — Ramp AI Index, 12 Aug 2026.



The median firm spends about what a single streaming subscription costs. It is not that they have a worse model; everyone rents the same models. They have not found work worth pointing one at.

Takeaway The same technology, bought at 620x different intensity. Which use cases are worth it — and to whom — is C1.

AI in the news — compute becomes an asset class

THIS MONTH

The largest financing story of the month, and the half of the bill it does not touch.

US\$500bn, arranged for compute. On 10 August NVIDIA announced financing platforms with **Apollo, BlackRock, Blackstone, Brookfield, Goldman Sachs and KKR**, to mobilise over **US\$500bn** of third-party capital.

Why it can be financed at all: Huang's own argument — NVIDIA compute is “fungible and transferable across customers and operators”. It is collateral because somebody else can use it.

What that money does not finance. Your data pipeline. Your eval harness. Your integration into the systems the work actually lives in. Nobody will lend against those, for the same reason: they are worth nothing to anyone else.

Why today: C1 asks what to own and what to rent. C2 shows this capital coming back to you inside the price of a token — someone else's fixed cost, resold as your variable cost.

The named borrowers are frontier labs, enterprises and AI cloud operators. Source: NVIDIA newsroom, 10 Aug 2026, reported the same day by CNBC; contingent on final agreements.

Takeaway Compute is becoming rentable infrastructure. The half that decides your unit economics is not.

What you read before today

Five briefs — four foundations plus the case brief. Each one seeds a capsule. Here is the through-line.

#	The reading	Where it lands
1	Iansiti & Lakhani — Competing in the Age of AI — the AI factory, scale / scope / learning, network effects, and the operating model that carries them	C1 • C2 • C3
2	McKinsey — The State of AI — adoption is finished and value is not; and how to read a number before you repeat it	C1 • C2 • C3
3	AI unit economics — cost to serve — the specification for the calculator you build after the break; every number in C2 comes from here	C2 (and C1 • C3)
4	GraphCast & Climate TRACE — what a vertical needs in place before AI is even possible	C1 • C3
5	Case brief — Ant Group (A) — the company Reading 1 makes its exhibit; 22 pages that do not flatter it	C3 — the case

Reading 1 is the framework and Ant is the framework's own exhibit — which is why the case brief travels with it rather than standing alone.

Takeaway Everything you read gets used today — five briefs, three capsules, one company.

Class 5 — objective & learning outcomes

FROM THE OBTL+
OUTLINE • WEEK 5

Straight from the course outline, so you can see the line from today to assessment.

	Upon successful completion of this course, you will be able to...	Today
ILO 1	Distinguish emerging AI technologies and select appropriate techniques for given business problem classes.	C3 — the two archetypes
ILO 2	Build and validate an end-to-end AI solution using both low-code platforms and a code-based stack.	C2 — the calculator → A2
ILO 3	Evaluate design alternatives against business factors (cost, time-to-deploy, risk, ROI) and technical factors, and justify the chosen approach.	C1 • C2 — primary
ILO 4	Communicate an AI solution — its problem framing, design choices, limitations and risks — to a mixed technical and business audience.	C3 — the case, primary

Week 5 in the outline: “AI strategy, economics, use-case selection — use-case framework; data readiness; build vs buy at scale; AI factory + network effects; cost-to-serve + ROI; vertical applications.”

Takeaway Today: choose the use case, price it honestly, and defend the choice to someone who controls a budget.

The guest slot — what happened, and what we did

You were told an industry guest would close today. He cannot make it. Here is the honest version.

We did not drop what he was going to cover. We asked him for the six things he intended to teach, and we are teaching all six ourselves. They are badged through the deck so you can see exactly which slides they are.

Look for: a navy outline chip reading **INSTEAD OF THE GUEST LECTURE** — and, where the slide carries my own experience of making the call, a second chip reading **FROM PRACTICE**.

This is also today's argument, happening to us. A plan met a constraint. The choice was to abandon the outcome or to re-cost the way of getting it. We re-costed: six topics absorbed, no teaching cut, two activities trimmed by two minutes each.

The honest cost: the practitioner war stories are mine rather than his, and there are fewer of them. You should know which parts of a session were bought and which were built.

Six absorbed topics: the kill story • the intake screen • economics at scale • the long tail of cost-to-serve • pilot to product • judgement as the product. The next slide places each one.

Takeaway Nothing on the syllabus was dropped. The delivery was re-costed — which is the week's own lesson.

The six, and where each one lands

Every topic the guest was briefed to cover, and the question it answers when we reach it.

	The slide	The question it answers
C1	One we killed — and the tell we missed	What did we see before the money was spent, and choose not to act on?
C1	The intake screen	What evidence buys a pilot, what buys a budget line, and who is allowed to say no?
C2	When the economics only work at scale	Why is the same use case obviously right for one firm and absurd for another?
C2	What cost-to-serve really contains	What is in the bill once it leaves the slide — retries, review, integration, the engineer?
C2	From pilot to product	What changes in the arithmetic the day a capability acquires customers?
C3	When the model is a commodity, the judgement is the product	If everyone rents the same models, what is left that is yours?

Takeaway Six topics, six slides, all three capsules — badged, so you can hold us to it.

A1 has closed — and A2 is running

Ninety seconds on assessment, because one deadline this week is easy to miss.

A1 Foundations Build closed on Sunday 30 August. You each get **written feedback against the rubric**, through NTULearn. There is no cohort-wide mark-distribution session — 175 submissions cannot be marked in a day, and I would rather your own feedback be good than a class summary be fast. **Nothing to do today:** if you submitted, you are done. If something went wrong with the submission itself, tell me in the break rather than by email at midnight.

A2 Applied AI System — the date people miss is Friday. Team declaration: Friday 4 September, 23:59 SGT — a file in the submission folder, not an email. Teams are your existing course groups; they are not re-formed. Starter material Wednesday 2 September. Due Sunday 13 September; peer rating Wednesday 16 September.

The non-contribution rule, said once: peer rating is a **signal, not a verdict**. A non-contribution finding needs **two corroborating signals** — the peer rating and no traceable work in the commit history or CONTRIBUTIONS.md. It draws a written warning first, and a confirmed finding scores zero.

All deadlines 23:59 SGT. A2 is released; the brief, the FAQ and the reference datasets are on NTULearn.

Takeaway One date to write down today: **Friday 4 September, 23:59 — the team declaration.**

Today's plan

Where each piece happens, and roughly how long. One 15-minute break, the same for all three cohorts.

	Block	What happens	Min
C0	Where we are • what's new • today's plan	recap • news • quiz • the guest • A1 & A2	~8
C1	Use-case selection & the AI factory	lecture • 2 polls • 2 quick checks • group activity	~49
C2	What does it actually cost? — Part A	lecture • poll • the worked example	~35
—	Break		15
C2	...continued — you build the calculator	notebook • pair exercise on a new scenario	~28
C3	Two archetypes + the Ant Group case	case brief • discussion • share-out	~46
—	Close — takeaways and what is coming in Class 6		~8

Takeaway Three capsules: which use cases are worth it → what they cost → what that means as strategy.