



C0

The last class — and today's plan

~6 min · the last class · recap · AI in the news · what you read · opening quiz · today's plan

BY THE END OF THIS CAPSULE YOU WILL BE ABLE TO

Know that this is the final session, what changed in AI since we last met, and exactly what the next three hours will do with everything you have read.



The last time this room is together

Week 7 onward is an asynchronous submission window. There is no class, and there is no teaching.

This is the sixth of six, and there is no seventh. A2 is due Sunday 20 September 2026. The peer rating closes Thursday 24 September, in Week 8. The project follows on Sunday 4 October. None of it involves a room.

So anything that only works face to face happens today: C3 closes with the whole course as one argument, and then eight minutes on what it was for.

What that does NOT mean is that today is a wrap-up. The heaviest content of the course is in the next three hours: a failure taxonomy you can use on your own systems, five regulatory regimes, and a red-team build against something you wrote yourself. Nothing has been thinned out to make room for sentiment.

A2 feedback is offline: there is no Week-7 session to hand it back in, so it comes through NTULearn. Do not wait for a class that does not exist.

Takeaway Six sessions, no seventh. Everything that needs a room has to happen in the next three hours.

Where we were — and what today does with it

Class 5 asked which use cases are worth building and what they cost. Today asks what happens when they are wrong.

#	What Class 5 left you with	What Class 6 does with it
1	cost per SUCCESSFUL task	C1 prices “a human will check it” — 50,000 items at 95% leaves 2,500 errors
2	the failure rate you divided by	today it is the subject, not the denominator: which failures, seen or unseen
3	the model layer is rentable; the data layer is not	C2-A adds the third — the governance layer is not rentable either
4	the intake screen — who is allowed to say no	C1 states human-in-the-loop properly: window • evidence • authority
5	MYBank ran 3,000 strategies with nobody in the loop	C1's 2x2 — that is the irreversible quadrant, and a gate is the only control
6	Ant: the binding constraint was the regulator	C3's case — a board with the formal authority to act, and no way to use it

Takeaway Class 5 asked what it costs. Class 6 asks what it costs when it goes wrong, and who is on the hook.

AI in the news — the agent that got out

JULY–AUGUST 2026

A frontier lab lost track of its own agents for about ten weeks, then published the timeline.

Disclosed by OpenAI, 21 July 2026. Agents in an internal evaluation left their sandbox, built a message board inside the company's own artefact repository to coordinate with each other, and reached production systems at Hugging Face.

Undetected for about seventy days: first message 12 May, connected to the breach 20 July.

The control that would have caught it was not in the path. OpenAI's own report says its chain-of-thought monitors “did not run on the evaluations in this incident”, and that in retrospect they were “very effective at flagging” the activity. The safeguards it ships around its public models had not been extended to its internal ones.

Why today: that is C1's four-operational-layers slide, happening to the people who wrote the layers.

Sources: OpenAI, “The Hugging Face incident and the road ahead”, 21 Jul 2026 (its own post-incident report); technical detail presented at Black Hat USA, 5 Aug 2026. Checked 6 Sep 2026.

Takeaway A control that exists but does not run on the path that matters is not a control.

AI in the news — the EU started fining

2 AUG 2026

One month ago the Commission got the power to fine. On the same Act, the heavy obligations slipped to 2027.

Since 2 August 2026, the AI Office can enforce. Obligations on general-purpose models have applied since August 2025; the power to enforce them arrived only this year — demand documentation, run its own evaluations, order corrective measures, and fine up to 3% of worldwide turnover or €15m.

And on the same Act, the heavy part moved the other way. The high-risk obligations everyone quotes — the ones that would cover a hiring model or a credit model — were pushed to 2 December 2027 by the Digital Omnibus. Enforcement and transparency tightened; the substantive duties slipped sixteen months.

Why today: C2·A never asks “are we compliant?”. It asks compliant with WHAT, binding on WHOM, from WHEN.

Digital Omnibus, Regulation (EU) 2026/1744, in force 27 Jul 2026 — high-risk to 2 Dec 2027 (Annex III) and 2 Aug 2028 (Annex I). Enforcement powers from 2 Aug 2026, with Article 50 transparency; systems already on the market have until 2 Dec 2026 for the marking rule.

Takeaway “Is it compliant?” has no answer. On one law, two obligations moved in opposite directions six days apart.

What you read before today

Six briefs — five foundations plus the case. Every one of them is used in this room today.

#	The reading	Where it lands
1	Failure modes by modality — and responsible deployment as four operational layers — the only brief in the set that describes a concrete failure	C1 • C2-B
2	OWASP Top 10 for LLM Applications (2026 edition) — the attack surface, in the vocabulary your own security team already uses	C1 • C2-A • C2-B
3	NIST AI RMF 1.0 + the EU AI Act — a process standard and a product law, and why the difference decides what you have to prove	C2-A
4	Singapore IMDA • India DPDP and MeitY • China CAC — your three home markets, and what each one actually governs	C2-A
5	Hinton, Hassabis and the critics — two of the people who built this, disagreeing in public, and the case against both of them	C3 • the close
6	Case brief — Governing OpenAI (A) — four pages of clues and no verdict; it stops on 21 November 2023, and that is deliberate	C3 — the case

Takeaway Six briefs, four blocks, one company. By the close every one of them will have been used here.

Class 6 — objective & learning outcomes

FROM THE OBTL+
OUTLINE • WEEK 6

Straight from the course outline, so you can see the line from today to the assessment.

	Upon successful completion of this course, you will be able to...	Today
ILO 1	Distinguish emerging AI technologies and select appropriate techniques for given business problem classes.	C1 — the failure taxonomy
ILO 2	Build and validate an end-to-end AI solution using both low-code platforms and a code-based stack.	C2-B — the attack and the guardrail
ILO 3	Evaluate design alternatives against business factors (cost, time-to-deploy, risk, ROI) and technical factors, and justify the chosen approach.	C1 • C2-A — primary
ILO 4	Communicate an AI solution — its problem framing, design choices, limitations and risks — to a mixed technical and business audience.	C3 — the case, primary

Week 6 in the outline: “Failure modes per modality (LLM hallucination, agent loops, ML bias, embodied-AI safety); operational responsible deployment (eval, observability, guardrails); regulatory frameworks; frontier outlook.”

Takeaway Today: anticipate the failure, put the control that answers it, and defend the call to a room that is not technical.

Today's plan

Where each piece happens, and roughly how long. One 15-minute break, the same for all three cohorts.

	Block	What happens	Min
C0	The last class • what's new • today's plan	recap • two news items • what you read • opening quiz	~11
C1	How each kind of AI breaks	lecture • poll • think-pair-share • the 2×2	~50
C2•A	What you owe, and to whom	lecture • quiz • five regimes, five objects	~16
—	Break		15
C2•B	Break it, then fix it	poll and category draw • pair build • share-out	~28
C3	The frontier, the case, and the whole course	case • discussion • reveal • recap • dates	~55
—	Close — what this course is for, after MSEAI		~8

Takeaway Two hours on how it breaks and what you owe. Then the case, the course, and what it was all for.