

CLOSE

PE6201 — Emerging AI Technologies

~8 min · the last thing you hear in this room · no logistics, they finished in C3

THE CLOSE

Six weeks, one argument — and what it is for once you leave this room.



Where this course sits in your AI journey

It was deliberately the broad one. That was a design decision, not a limitation.

You can now use these words correctly: token • embedding • chunk • context window • system prompt vs user prompt • structured output • retrieval • eval set • tool contract • step cap • drift • cost per successful task.

In a vendor meeting that vocabulary is the difference: between being briefed and being sold to.

And it came from building, not from slides. Tokenisation and embeddings • chunking and retrieval • prompting and structured output • eval harness design • agent loops and step caps • unit economics • drift and guardrails. You did NOT train a model, tune serving, or run anything at scale.

Wherever you go deeper — inside MSEAI or anywhere else: every specialist course begins from a method somebody already chose. Choosing it is the part nobody teaches twice.

Takeaway The broad one is what tells you which deep one you actually need — wherever you go and take it.

The portfolio you have not noticed you have

Nine things you built, or are about to. Most of them are already on GitHub.

What you built	When	What it demonstrates to someone hiring
A discriminative classifier benchmarked against a general-purpose API	1	That you cost and test the cheap option before reaching for the expensive one
A tokeniser walkthrough — what the model actually reads	1	That you know text becomes numbers before it becomes an answer
A RAG pipeline over your own documents	2	That you understand retrieval as an architecture, not as a product name
A structured-output prompt plus a ten-case evaluation harness	3	That you define a measurement before you claim an improvement
An agent with five tools — and its failure trace	4	That you have seen reliability compound downward and can explain why
A cost-to-serve calculator for a real use case	5	That you can put a number on an AI proposal, which most people cannot
A working red-team attack and the guardrail that closed it	6 • today	That you have broken something on purpose and fixed it in code
A2 — an agent with tools, guardrails, an eval set and a cost model	A2	A whole system with its own evidence — including two failures you built on purpose
The project — repo, recorded demo, and the trade-off argument	Project	The one you will actually show someone

Takeaway That is a portfolio. Most people finishing a course like this cannot name one artefact.

Where this maps, honestly

Seven recognisable job titles. The eighth has no title at all, and it is on the next slide.

The role	What they actually do	Built on
AI PRODUCT MANAGER	Chooses the method, sets the measurement, owns the cost per successful task.	W1 • W3 • W5
AI / ML ENGINEER	Builds and ships it — retrieval, the tool layer, evals, guardrails, and the bill.	W2 • W3 • W4 • A2
SOLUTION ARCHITECT	Decides what to own and what to rent, and how the pieces fit together.	W2 • W4 • W5 • A2
DATA / PLATFORM ENGINEER	Owns the one layer nobody can rent — pipelines, the corpus, and who may write into it.	W2 • W5 • W6
AI RISK, ASSURANCE & GOVERNANCE	Writes and tests the controls. Regulated sectors are hiring now and cannot fill it.	W6
PROGRAMME / TRANSFORMATION LEAD	Sequences the deployment, owns the change, decides what a pilot has to prove.	Moderna • DBS • AI Wars • Anthropic
AI STRATEGY & ADVISORY • FOUNDER	Advises on build-versus-buy and vendor choice — or bets a company on the answer.	W1 • W5 • Anthropic • Ant • A2 • the project

Takeaway Seven titles. The eighth has no title, and it is the one most of this room should hear about.

The role with no job title

It is the one about two thirds of this room is best placed for, and nobody advertises it.

Every team ends up with a translator. They sit in finance, legal, clinical or operations — not in engineering — and they are the one person who can tell whether the vendor is lying and what the thing will cost at volume.

You did not need to become technical for this. You needed to know what questions survive contact with a confident answer.

That person is scarce because the two halves rarely meet. Technical people who understand the deployment context are rare; domain people who understand what the model is doing are rarer. Building the small version yourself puts you in the second.

Which is why: you built a classifier, a RAG system and an agent, and you priced them per successful task — not to become an engineer, but so that nobody can tell you a story about one.

Takeaway You do not have to be the person who builds it to be the most useful person in the conversation.

Six habits that outlive every model version

Every price, ranking and capability claim in this course will be wrong within a year. These will not.

COST PER SUCCESSFUL TASK

Not per token. If a third of outputs need redoing, your real unit cost is twice what the invoice says.

From: Week 5

YOUR TEN CASES BEAT THEIR BENCHMARK

A public score measures their task. An eval set you wrote measures yours — and you write it before you tune anything.

From: Week 3

A PROMPT ASKS. A PERMISSION ENFORCES.

Eleven capitalised instructions not to touch production did not stop the agent deleting the database. If it matters, it lives in code.

From: Week 4

READ WHAT THE BENCHMARK MEASURED

Before the number, the task definition. Most capability arguments dissolve once you read the benchmark's own paper.

From: Weeks 3 • 5 • 6

WHO WAS ASKED, AND WHAT COUNTED

Who was asked, what threshold counted as success, and what the denominator was. Three questions, any field.

From: Weeks 5 & 6

GATE THE IRREVERSIBLE, MONITOR THE INVISIBLE

In that order. The two-question diagnostic, and the whole of responsible deployment in six words.

From: Week 6

Takeaway Carry the method. The facts have a shelf life; the questions do not.

THE LAST WORD

You will be among the people who find out.

Two of the people who built this technology disagree, in public, about whether it is dangerous — and neither of them is being careless. That argument is not settled, and it will not be settled by an expert. It will be settled by what gets deployed, by whom, with what gate on it. Which is your job now. Thank you — it has been a genuine pleasure.

Ajay Vikram Singh

ajay.ai.education@gmail.com · +65 9189 6093 (phone · WhatsApp · FaceTime)